


Tidyverse

Collection of R packages

What is *tidyverse* ?

- The tidyverse is a collection of R packages designed for data science.
- All packages share an underlying design philosophy, grammar, and data structures.
- Developed by Hadley Wickham



Components of *tidyverse*

- **ggplot2:** ggplot2 is a system for declaratively creating graphics, based on The Grammar of Graphics.
- **dplyr:** dplyr provides a grammar of data manipulation, providing a consistent set of verbs that solve the most common data manipulation challenges
- **tidyr:** tidyr provides a set of functions that help you get to tidy data.
- **readr:** readr provides a fast and friendly way to read rectangular data (like csv, tsv, and fwf).
- **purrr:** purrr enhances R's functional programming (FP) toolkit by providing a complete and consistent set of tools for working with functions and vectors.
- **tibble:** tibble is a modern re-imagining of the data frame, keeping what time has proven to be effective, and throwing out what it has not.
- **stringr:** stringr provides a cohesive set of functions designed to make working with strings as easy as possible
- **forcats:** forcats provides a suite of useful tools that solve common problems with factors

Loading *tidyverse*

- All packages in tidyverse can be installed and loaded at one go

```
install.packages("tidyverse")
```

```
library(tidyverse)
```

Package *dplyr*

Handling the Data Efficiently

class tbl_df

- We can create an object of class `tbl_df`. It can be created by function `as_tibble`

Syntax : `as_tibble(objDF)`

where

`objDF`: An object of class `data.frame`, or a list with each element with same length

tibble Object

- A tibble is a modern class of data frame within R
- It has a convenient print method, will not convert strings to factors, and does not use row names

```
> dd = as_tibble(mtcars)
> class(dd)
[1] "tbl_df"      "tbl"        "data.frame"
> dd
# A tibble: 32 x 11
   mpg   cyl  disp    hp  drat    wt   qsec    vs
*   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1    21     6   160   110   3.9   2.62  16.5     0
2    21     6   160   110   3.9   2.88  17.0     0
3   22.8     4   108    93   3.85   2.32  18.6     1
4   21.4     6   258   110   3.08   3.22  19.4     1
5   18.7     8   360   175   3.15   3.44  17.0     0
6   18.1     6   225   105   2.76   3.46  20.2     1
7   14.3     8   360   245   3.21   3.57  15.8     0
8   24.4     4   147.    62   3.69   3.19   20      1
9   22.8     4   141.    95   3.92   3.15  22.9     1
10  19.2     6   168.   123   3.92   3.44  18.3     1
# ... with 22 more rows, and 3 more variables:
#   am <dbl>, gear <dbl>, carb <dbl>
```

Functions in **dplyr**

- **arrange**: reordering rows in the data frame
- **select**: selecting columns / variables
- **filter**: selecting rows / observations
- **rename**: renaming variables
- **mutate**: adding new columns to the data frame
- **summarize** / **summarise**: generating summary statistics of the data frames

Arranging the rows

- The rows in `tbl_df` object can be arranged using function `arrange`

Syntax : `arrange(Obj_tbl_df, col1,col2,...)`

Where

`Obj_tbl_df` : `tbl_df` object

`col1, col2,...` : Columns for sorting the data

Arrange Example

```
> tbl_Cars
# A tibble: 93 x 27
  Manufacturer Model      Type Min.Price Price Max.Price MPG.city MPG.highway
*   <fctr>      <fctr>    <fctr>   <dbl>  <dbl>   <dbl>   <int>    <int>
1     Acura   Integra   Small    12.9   15.9    18.8     25      31
2     Acura   Legend  Midsize    29.2   33.9    38.7     18      25
3       Audi     90  Compact    25.9   29.1    32.3     20      26
4       Audi    100  Midsize    30.8   37.7    44.6     19      26
5       BMW    535i  Midsize    23.7   30.0    36.2     22      30
6     Buick   Century  Midsize    14.2   15.7    17.3     22      31
7     Buick   LeSabre   Large    19.9   20.8    21.7     19      28
8     Buick Roadmaster  Large    22.6   23.7    24.9     16      25
9     Buick   Riviera  Midsize    26.3   26.3    26.3     19      27
10    Cadillac DeVille   Large    33.0   34.7    36.3     16      25
# ... with 83 more rows, and 19 more variables: AirBags <fctr>, DriveTrain <fctr>,
# Cylinders <fctr>, EngineSize <dbl>, Horsepower <int>, RPM <int>, Rev.per.mile <int>,
# Man.trans.avail <fctr>, Fuel.tank.capacity <dbl>, Passengers <int>, Length <int>,
# Wheelbase <int>, Width <int>, Turn.circle <int>, Rear.seat.room <dbl>,
# Luggage.room <int>, Weight <int>, Origin <fctr>, Make <fctr>
```

```
> ord_Model <- arrange(tbl_Cars , Model)
> ord_Model
# A tibble: 93 x 27
  Manufacturer Model      Type Min.Price Price Max.Price MPG.city MPG.highway
  <fctr>      <fctr>    <fctr>   <dbl>  <dbl>   <dbl>   <int>    <int>
1       Audi    100  Midsize    30.8   37.7    44.6     19      26
2 Mercedes-Benz 190E  Compact    29.0   31.9    34.9     20      29
3       Volvo   240  Compact    21.8   22.7    23.5     21      28
4 Mercedes-Benz 300E  Midsize    43.8   61.9    80.0     19      25
5       Mazda   323   Small     7.4    8.3     9.1     29      37
6       BMW    535i  Midsize    23.7   30.0    36.2     22      30
7       Mazda   626  Compact    14.3   16.5    18.7     26      34
8       Volvo   850  Midsize    24.8   26.7    28.5     20      28
9       Audi     90  Compact    25.9   29.1    32.3     20      26
10      Saab    900  Compact    20.3   28.7    37.1     20      26
# ... with 83 more rows, and 19 more variables: AirBags <fctr>, DriveTrain <fctr>,
# Cylinders <fctr>, EngineSize <dbl>, Horsepower <int>, RPM <int>, Rev.per.mile <int>,
# Man.trans.avail <fctr>, Fuel.tank.capacity <dbl>, Passengers <int>, Length <int>,
# Wheelbase <int>, Width <int>, Turn.circle <int>, Rear.seat.room <dbl>,
# Luggage.room <int>, Weight <int>, Origin <fctr>, Make <fctr>
```

Arranging Multiple Columns Example

```
> ord_mnf_md1 <- arrange(tbl_Cars,Manufacturer,Model)
> ord_mnf_md1
# A tibble: 93 x 27
  Manufacturer      Model      Type Min.Price Price Max.Price MPG.city MPG.highway
  <fctr>          <fctr> <fctr>   <dbl> <dbl>   <dbl>   <int>   <int>
1      Acura      Integra  Small    12.9  15.9    18.8     25     31
2      Acura      Legend  Midsize   29.2  33.9    38.7     18     25
3       Audi       100  Midsize   30.8  37.7    44.6     19     26
4       Audi       90  Compact   25.9  29.1    32.3     20     26
5       BMW       535i  Midsize   23.7  30.0    36.2     22     30
6      Buick      Century  Midsize   14.2  15.7    17.3     22     31
7      Buick     LeSabre   Large   19.9  20.8    21.7     19     28
8      Buick     Riviera  Midsize   26.3  26.3    26.3     19     27
9      Buick Roadmaster   Large   22.6  23.7    24.9     16     25
10     Cadillac  DeVille   Large   33.0  34.7    36.3     16     25
# ... with 83 more rows, and 19 more variables: AirBags <fctr>, DriveTrain <fctr>,
# Cylinders <fctr>, EngineSize <dbl>, Horsepower <int>, RPM <int>, Rev.per.mile <int>,
# Man.trans.avail <fctr>, Fuel.tank.capacity <dbl>, Passengers <int>, Length <int>,
# Wheelbase <int>, Width <int>, Turn.circle <int>, Rear.seat.room <dbl>,
# Luggage.room <int>, Weight <int>, Origin <fctr>, Make <fctr>
```

Arranging Multiple Columns Example

```
> ord_mnf_md1 <- arrange(tbl_Cars,Manufacturer,desc(Model))
> ord_mnf_md1
# A tibble: 93 x 27
  Manufacturer      Model      Type Min.Price Price Max.Price MPG.city MPG.highway
  <fctr>          <fctr>    <fctr>   <dbl>  <dbl>    <dbl>   <int>    <int>
1      Acura      Legend Midsize    29.2   33.9     38.7     18      25
2      Acura    Integra  Small    12.9   15.9     18.8     25      31
3       Audi        90 Compact    25.9   29.1     32.3     20      26
4       Audi       100 Midsize    30.8   37.7     44.6     19      26
5        BMW      535i Midsize    23.7   30.0     36.2     22      30
6     Buick Roadmaster  Large    22.6   23.7     24.9     16      25
7     Buick   Riviera Midsize    26.3   26.3     26.3     19      27
8     Buick   LeSabre  Large    19.9   20.8     21.7     19      28
9     Buick   Century Midsize    14.2   15.7     17.3     22      31
10    Cadillac Seville Midsize    37.5   40.1     42.7     16      25
# ... with 83 more rows, and 19 more variables: AirBags <fctr>, DriveTrain <fctr>,
# Cylinders <fctr>, EngineSize <dbl>, Horsepower <int>, RPM <int>, Rev.per.mile <int>,
# Man.trans.avail <fctr>, Fuel.tank.capacity <dbl>, Passengers <int>, Length <int>,
# Wheelbase <int>, Width <int>, Turn.circle <int>, Rear.seat.room <dbl>,
# Luggage.room <int>, Weight <int>, Origin <fctr>, Make <fctr>
```

Selecting Columns

- For selecting specific columns, we can use select function

Syntax : `select(objtbl , col 1, col 2, ...)`

OR

`select(objtbl , col 1:col n)`

where

objtbl: Object of class `tbl_df`

Select Examples

```
> select(tbl_Cars, 1:3)
# A tibble: 93 x 3
  Manufacturer Model      Type
*   <fctr>      <fctr>   <fctr>
1     Acura   Integra   Small
2     Acura   Legend   Midsize
3     Audi     90      Compact
4     Audi    100      Midsize
5     BMW     535i     Midsize
6     Buick   Century   Midsize
7     Buick   LeSabre   Large
8     Buick   Roadmaster Large
9     Buick   Riviera   Midsize
10    Cadillac DeVille   Large
# ... with 83 more rows
```

```
> select(tbl_Cars, Model:Max.Price)
# A tibble: 93 x 5
  Model      Type Min.Price Price Max.Price
*   <fctr>   <fctr>   <dbl> <dbl>   <dbl>
1   Integra Small     12.9  15.9    18.8
2   Legend  Midsize    29.2  33.9    38.7
3     90    Compact    25.9  29.1    32.3
4    100    Midsize    30.8  37.7    44.6
5   535i    Midsize    23.7  30.0    36.2
6   Century Midsize    14.2  15.7    17.3
7   LeSabre Large     19.9  20.8    21.7
8 Roadmaster Large     22.6  23.7    24.9
9   Riviera Midsize    26.3  26.3    26.3
10  DeVille  Large     33.0  34.7    36.3
# ... with 83 more rows
```

```
> select(tbl_Cars, ends_with("Price"))
# A tibble: 93 x 3
  Min.Price Price Max.Price
*   <dbl> <dbl>   <dbl>
1     12.9 15.9    18.8
2     29.2 33.9    38.7
3     25.9 29.1    32.3
4     30.8 37.7    44.6
5     23.7 30.0    36.2
6     14.2 15.7    17.3
7     19.9 20.8    21.7
8     22.6 23.7    24.9
9     26.3 26.3    26.3
10    33.0 34.7    36.3
# ... with 83 more rows
```

```
> select(tbl_Cars, starts_with("MPG"))
# A tibble: 93 x 2
  MPG.city MPG.highway
*   <int>      <int>
1      25         31
2      18         25
3      20         26
4      19         26
5      22         30
6      22         31
7      19         28
8      16         25
9      19         27
10     16         25
# ... with 83 more rows
```

Subsetting the data

- The data can be subsetting with function filter

Syntax : `filter(objtbl , criteria)`

where

objtbl: Object of class `tbl_df`

criteria: Condition of filtering

filter examples

```
> filter(tbl_Cars, Type=="Small")
# A tibble: 21 x 27
  Manufacturer Model   Type Min.Price Price Max.Price MPG.city MPG.highway AirBags DriveTrain
  <fctr>      <fctr> <fctr>   <dbl> <dbl>   <dbl>   <int>    <int>    <fctr>   <fctr>
1      Acura  Integra Small    12.9  15.9    18.8     25      31      None    Front
2      Dodge   Colt   Small     7.9   9.2    10.6     29      33      None    Front
3      Dodge  Shadow Small     8.4  11.3    14.2     23      29 Driver only Front
4      Eagle  Summit Small     7.9  12.2    16.5     29      33      None    Front
5      Ford  Festiva Small     6.9   7.4     7.9     31      33      None    Front
6      Ford  Escort  Small     8.4  10.1    11.9     23      30      None    Front
7       Geo   Metro  Small     6.7   8.4    10.0     46      50      None    Front
8      Honda   Civic Small     8.4  12.1    15.8     42      46 Driver only Front
9     Hyundai  Excel Small     6.8   8.0     9.2     29      33      None    Front
10    Hyundai Elantra Small     9.0  10.0    11.0     22      29      None    Front
# ... with 11 more rows, and 17 more variables: Cylinders <fctr>, EngineSize <dbl>, Horsepower <int>,
# RPM <int>, Rev.per.mile <int>, Man.trans.avail <fctr>, Fuel.tank.capacity <dbl>,
# Passengers <int>, Length <int>, Wheelbase <int>, Width <int>, Turn.circle <int>,
# Rear.seat.room <dbl>, Luggage.room <int>, Weight <int>, Origin <fctr>, Make <fctr>
```

```
> filter(tbl_Cars, Type=="Small" & Max.Price<10)
# A tibble: 6 x 27
  Manufacturer Model   Type Min.Price Price Max.Price MPG.city MPG.highway AirBags DriveTrain
  <fctr>      <fctr> <fctr>   <dbl> <dbl>   <dbl>   <int>    <int>    <fctr>   <fctr>
1      Ford  Festiva Small     6.9   7.4     7.9     31      33      None    Front
2     Hyundai  Excel Small     6.8   8.0     9.2     29      33      None    Front
3      Mazda   323   Small     7.4   8.3     9.1     29      37      None    Front
4     Pontiac  LeMans Small     8.2   9.0     9.9     31      41      None    Front
5      Subaru  Justy Small     7.3   8.4     9.5     33      37      None    4WD
6 Volkswagen  Fox   Small     8.7   9.1     9.5     25      33      None    Front
# ... with 17 more variables: Cylinders <fctr>, EngineSize <dbl>, Horsepower <int>, RPM <int>,
# Rev.per.mile <int>, Man.trans.avail <fctr>, Fuel.tank.capacity <dbl>, Passengers <int>,
# Length <int>, Wheelbase <int>, Width <int>, Turn.circle <int>, Rear.seat.room <dbl>,
# Luggage.room <int>, Weight <int>, Origin <fctr>, Make <fctr>
```


filter examples

```
> filter(tbl_Cars, Manufacturer %in% c("Acura","Audi"))
```

```
Source: local data frame [4 x 27]
```

	Manufacturer (fctr)	Model (fctr)	Type (fctr)	Min.Price (dbl)	Price (dbl)	Max.Price (dbl)	MPG.city (int)	MPG.highway (int)
1	Acura	Integra	Small	12.9	15.9	18.8	25	31
2	Acura	Legend	Midsize	29.2	33.9	38.7	18	25
3	Audi	90	Compact	25.9	29.1	32.3	20	26
4	Audi	100	Midsize	30.8	37.7	44.6	19	26

Variables not shown: AirBags (fctr), DriveTrain (fctr), cylinders (fctr), EngineSize (dbl), Horsepower (int), RPM (int), Rev.per.mile (int), Man.trans.avail (fctr), Fuel.tank.capacity (dbl), Passengers (int), Length (int), wheelbase (int), width (int), Turn.circle (int), Rear.seat.room (dbl), Luggage.room (int), weight (int), origin (fctr), Make (fctr)

Renaming Columns

- The columns can be renamed with function `rename()`

Syntax : `rename(objtbl, newname1=oldname1,
newname2=oldname2,...)`

where

`objtbl`: Object of class `tbl_df`

rename example

```
> rename(tbl_Cars, Minimum=Min.Price, Maximum=Max.Price)
# A tibble: 93 x 27
  Manufacturer Model Type Minimum Price Maximum MPG.city MPG.highway
*   <fctr>      <fctr> <fctr>   <dbl> <dbl>   <dbl>   <int>   <int>
1     Acura   Integra Small    12.9  15.9    18.8     25     31
2     Acura   Legend Midsize   29.2  33.9    38.7     18     25
3     Audi     90 Compact   25.9  29.1    32.3     20     26
4     Audi    100 Midsize   30.8  37.7    44.6     19     26
5     BMW     535i Midsize   23.7  30.0    36.2     22     30
6     Buick   Century Midsize   14.2  15.7    17.3     22     31
7     Buick   LeSabre Large    19.9  20.8    21.7     19     28
8     Buick Roadmaster Large    22.6  23.7    24.9     16     25
9     Buick   Riviera Midsize   26.3  26.3    26.3     19     27
10    Cadillac DeVille Large    33.0  34.7    36.3     16     25
# ... with 83 more rows, and 19 more variables: AirBags <fctr>, DriveTrain <fctr>,
# Cylinders <fctr>, EngineSize <dbl>, Horsepower <int>, RPM <int>, Rev.per.mile <int>,
# Man.trans.avail <fctr>, Fuel.tank.capacity <dbl>, Passengers <int>, Length <int>,
# Wheelbase <int>, Width <int>, Turn.circle <int>, Rear.seat.room <dbl>,
# Luggage.room <int>, Weight <int>, Origin <fctr>, Make <fctr>
```

Adding new column

- We can create one or more new columns / variables in the data with function mutate

Syntax : `mutate(objtbl , assign)`

where

objtbl: Object of class `tbl_df`

assign: Specification of assignment for new column

mutate example

```
tbl_Cars_rng <- mutate(tbl_Cars , Price_Range = Max.Price - Min.Price ,  
                        ratio = Weight/Passengers)
```

```
> select(tbl_Cars_rng,Model,Price_Range,ratio)  
# A tibble: 93 x 3  
   Model Price_Range ratio  
   <fctr>      <dbl>   <dbl>  
1  Integra      5.9 541.0000  
2  Legend      9.5 712.0000  
3    90       6.4 675.0000  
4   100     13.8 567.5000  
5  535i     12.5 910.0000  
6 Century      3.1 480.0000  
7 LeSabre      1.8 578.3333  
8 Roadmaster    2.3 684.1667  
9  Riviera      0.0 699.0000  
10 Deville      3.3 603.3333  
# ... with 83 more rows
```

Summarizing the data

- The data can be summarized with the function `summarize/summarise`

Syntax : `summarize(objtbl, assign)`

where

`objtbl`: Object of class `tbl_df`

`assign`: Specification of assignment for new column

summarize example

```
> summarize(tbl_Cars, avg_Price = mean(Price, na.rm = TRUE),  
+           sd_engSize = sd(EngineSize, na.rm = TRUE))  
Source: local data frame [1 x 2]
```

	avg_Price (dbl)	sd_engSize (dbl)
1	19.50968	1.037363

Grouping

- The `group_by` function takes an existing `tbl` and converts it into a grouped `tbl` where operations are performed "by group".

Syntax : `group_by(objtbl)`

where

`objtbl`: Object of class `tbl_df`

Group by example

```
> by_Air_Origin <- group_by(tbl_cars, Origin, AirBags)
> summarise(by_Air_Origin, avg_Price = mean(Price, na.rm = TRUE),
+           sd_engSize = sd(EngineSize, na.rm = TRUE))
Source: local data frame [6 x 4]
Groups: Origin [?]
```

	origin (fctr)	AirBags (fctr)	avg_Price (dbl)	sd_engSize (dbl)
1	USA	Driver & Passenger	24.57778	0.5600099
2	USA	Driver only	19.86957	1.2303796
3	USA	None	13.33125	0.9949874
4	non-USA	Driver & Passenger	33.24286	0.4270608
5	non-USA	Driver only	22.78000	0.7680974
6	non-USA	None	13.03333	0.5883676

Chaining / Pipelining

- We can pipeline the operations which are consecutive to one tbl object using %>% operator.

Syntax :

objtbl %>% operations

where

objtbl: Object of class tbl_df

Pipelining the data operations

```
##Instead of  
filter(select(tbl_Cars, Model, Price, Type) , Type=="Small")
```

```
## We can type  
tbl_cars %>%  
  select(Model, Price, Type) %>%  
  filter(Type=="Small")
```

```
#Considering  
x1 <- 1:5; x2 <- 2:6
```

```
##Instead of  
sqrt(sum((x1-x2)^2))
```


```
## We can type  
(x1-x2)^2 %>% sum() %>% sqrt()
```

Joins

- In package ***dplyr***, we have all the types of joins like
 - Inner join
 - Left Join
 - Right Join
 - Full Join

Inner Join

Intersection set




```
> A
  IdNum  A
1     1 234
2     2 134
3     3 145
4     4 653
5     5 246
```

```
> B
  IdNum  B
1     1 200
2     2 100
3     3 1444
4     6 400
5     7 160
```

```
> inner_join(A,B,by="IdNum")
  IdNum  A  B
1     1 234 200
2     2 134 100
3     3 145 1444
```

Left Outer Join



```
> A
```

	IdNum	A
1	1	234
2	2	134 ✓
3	3	145 ✓
4	4	653
5	5	246

```
> B
```

	IdNum	B
1	1	200
2	2	100
3	3	1444
4	6	400
5	7	160

```
> left_join(A,B,by="IdNum")
```

	IdNum	A	B
1	1	234	200
2	2	134	100
3	3	145	1444
4	4	653	NA
5	5	246	NA

Right Outer Join




```
> A
  IdNum  A
1      1 234
2      2 134
3      3 145
4      4 653
5      5 246
```

```
> B
  IdNum  B
1      1 200
2      2 100
3      3 1444
4      6 400
5      7 160
```

```
> right_join(A,B,by="IdNum")
  IdNum  A  B
1      1 234 200
2      2 134 100
3      3 145 1444
4      6 NA 400
5      7 NA 160
```

Full Outer Join



```
> A
  IdNum  A
1     1 234
2     2 134
3     3 145
4     4 653
5     5 246
```

```
> B
  IdNum  B
1     1 200
2     2 100
3     3 1444
4     6 400
5     7 160
```

```
> full_join(A,B,by="IdNum")
```

```
  IdNum  A  B
1     1 234 200
2     2 134 100
3     3 145 1444
4     4 653 NA
5     5 246 NA
6     6 NA 400
7     7 NA 160
```